

a year with
mongoDB

a talk by Armin '@mitsuhiko' Ronacher for PyGruun 2013

That's me.
I do Computers.
Currently at Fireteam / Splash Damage.
We do Internet for Pointy Shooty Games.



let's not beat
around the bush

I don't .(

like it .(

but we're not all so negative

“MongoDB is a pretty okay data store”

Jared Hefty (@bridwag)

this is not a rant
it's our experience in a nutshell
we find corner cases
draw your own conclusions

“MongoDB is like a nuclear reactor: ensure proper working conditions and it's perfectly safe and powerful.”

myself on 13th of October 2012

What changed?

RAD Soldiers

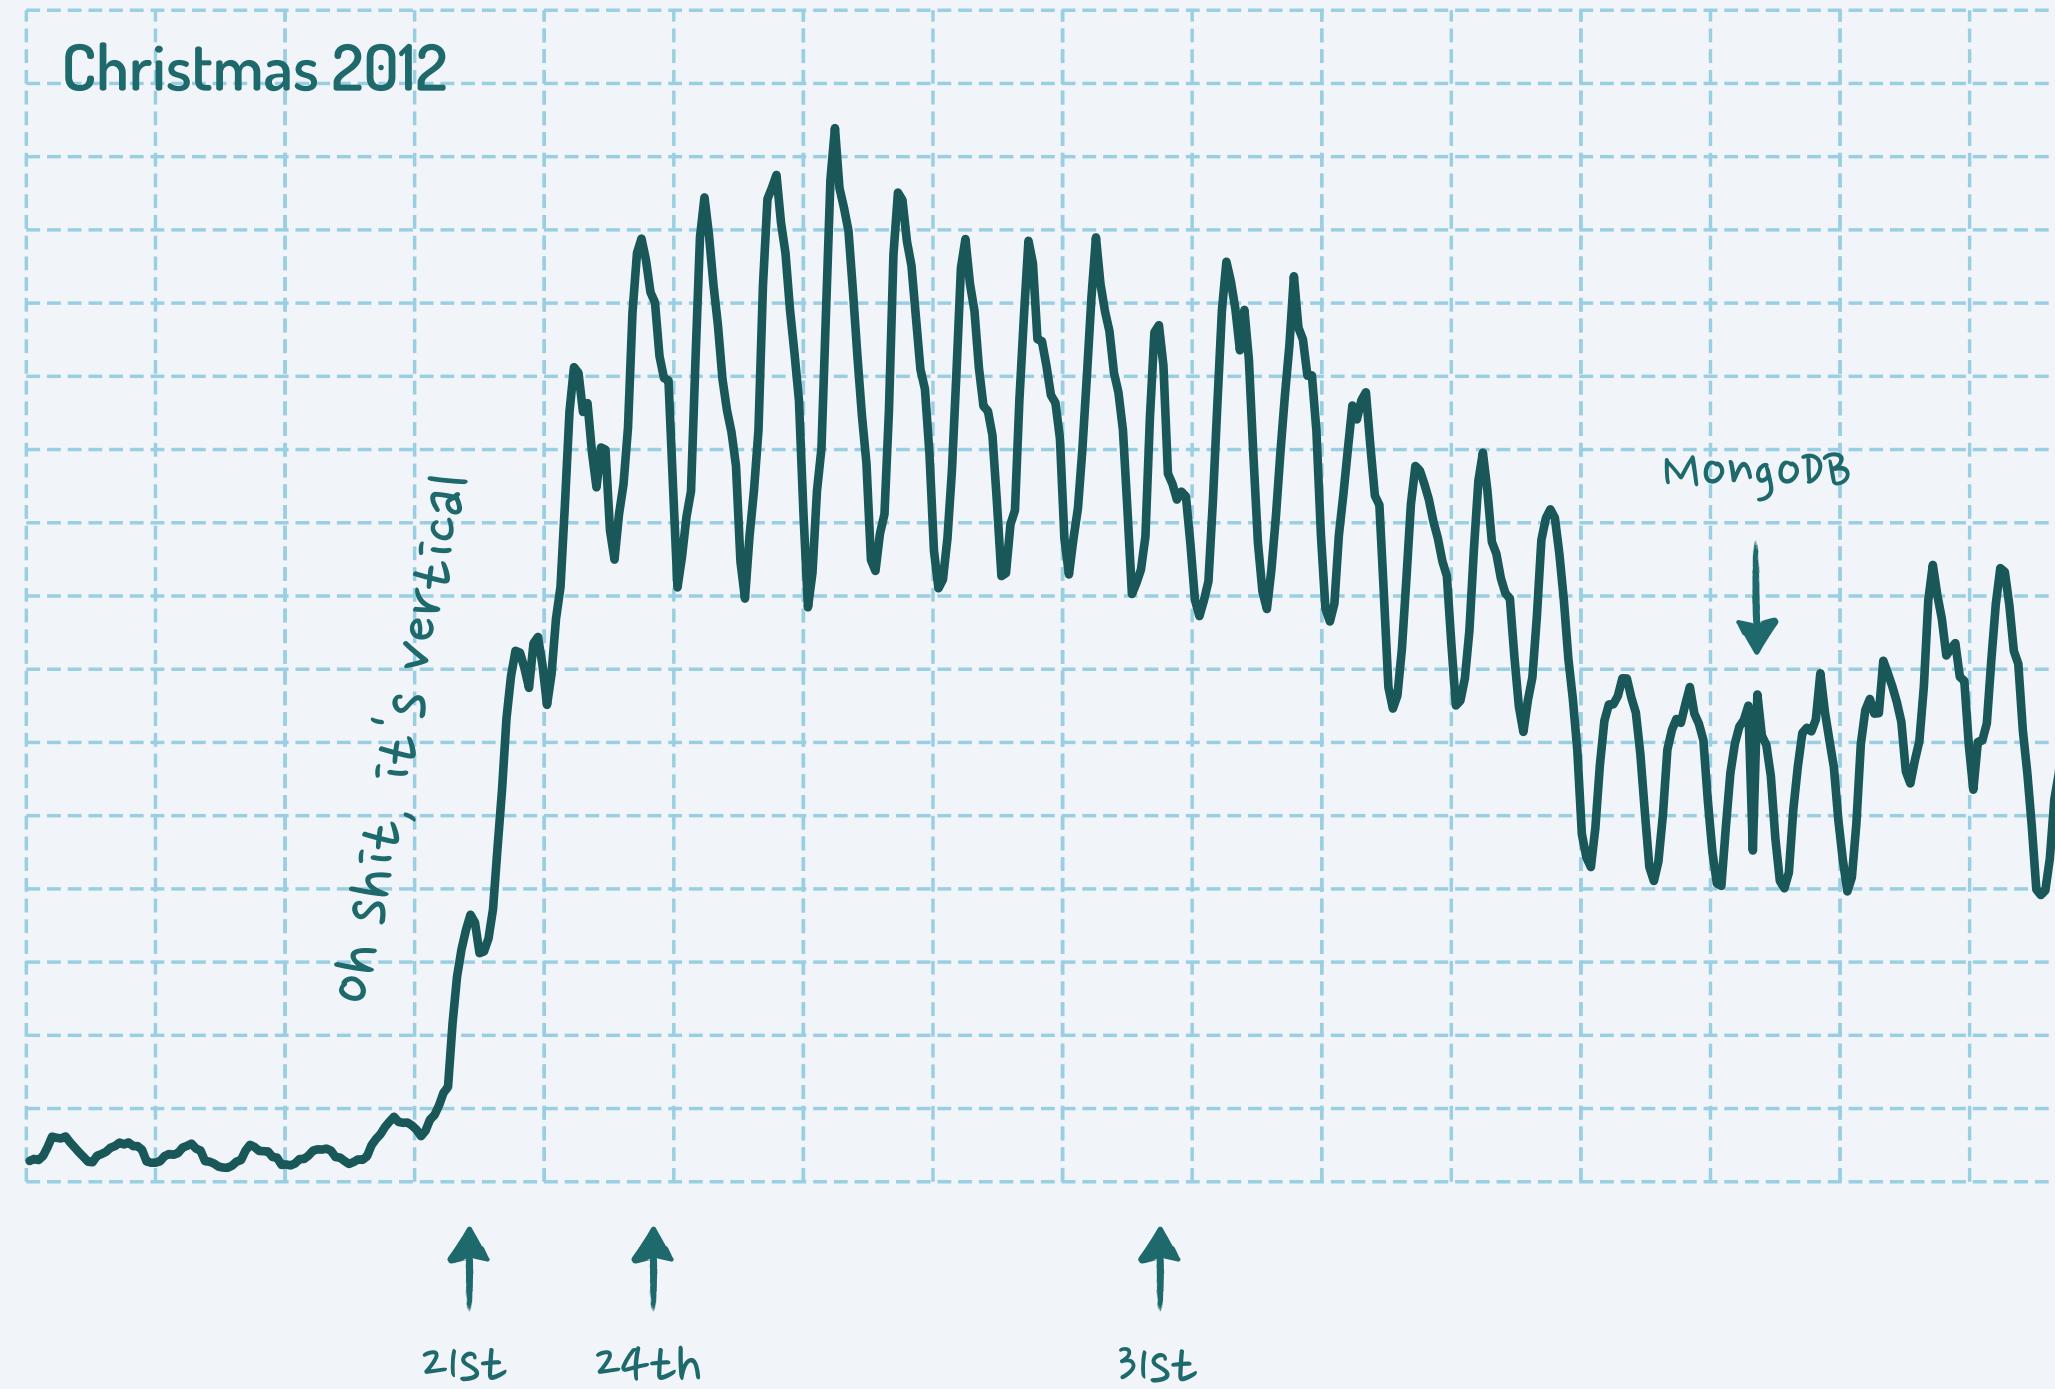


RAD
SOLDIERS

RAD Soldiers



RAD Soldiers API calls



?



{ MongoDB Overview }

We recently asked the question

WHY?

Why the fuck
did we pick
MongoDB?

schemaless Why the fuck
did we pick
MongoDB?

schemaless Why the fuck
scalable did we pick
MongoDB?

schemaless Why the fuck
scalable did we pick
simple MongoDB?

schemaless json records
scalable
simple

schemaless json records
scalable auto sharding
simple

schemaless json records
scalable auto sharding
simple think in records

schemaless is wrong
mongodb's sharding is annoying
thinking in records is hard
trololol: two-phase commit

mongod

mongoc

mongos

mongod mongods

mongoc

mongos

mongod mongods

mongoc mongocts

mongos

mongod mongods

mongoc mongocts

mongos mongoses

stores data

mongod mongods

mongoc mongocts

mongos mongoses

stores data mongod **mongods**

says what's where mongoc **mongoCS**

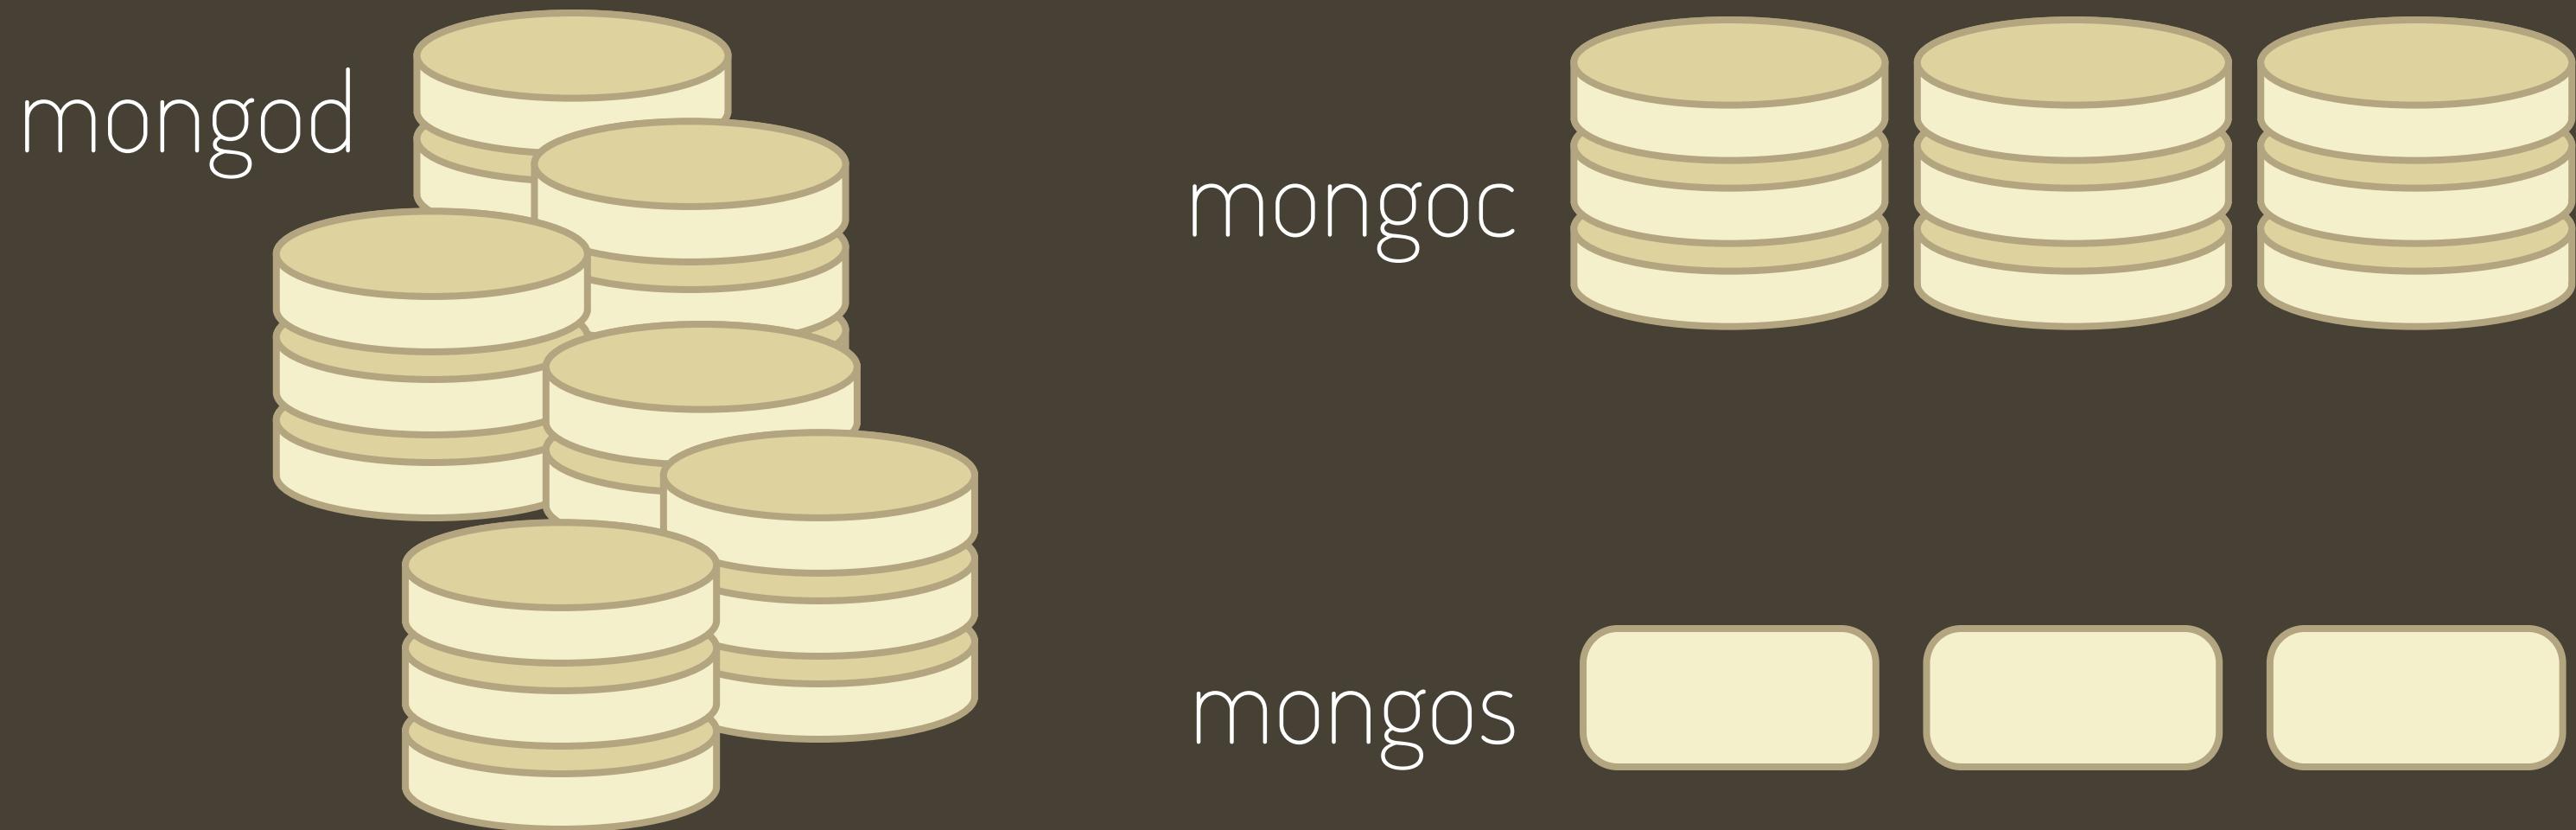
mongos **mongoses**

stores data mongod **mongods**

says what's where mongoc **mongoCS**

routes and merges mongos **mongoses**

Many Moving Parts



{ We Fail }

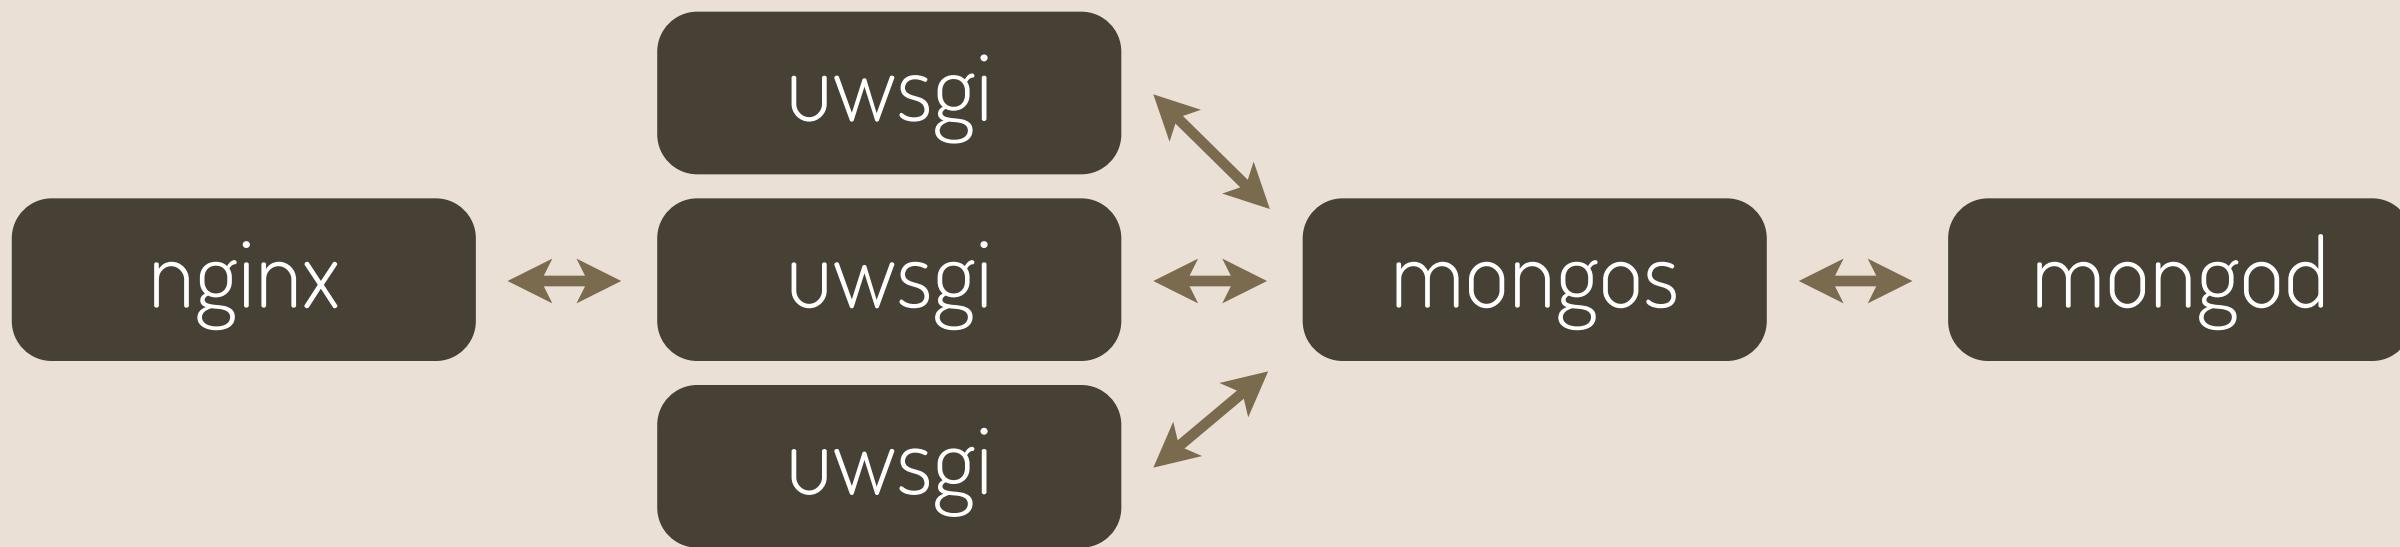
workers on m1.small
most of the time in IO wait
no need for more CPU

oh really?

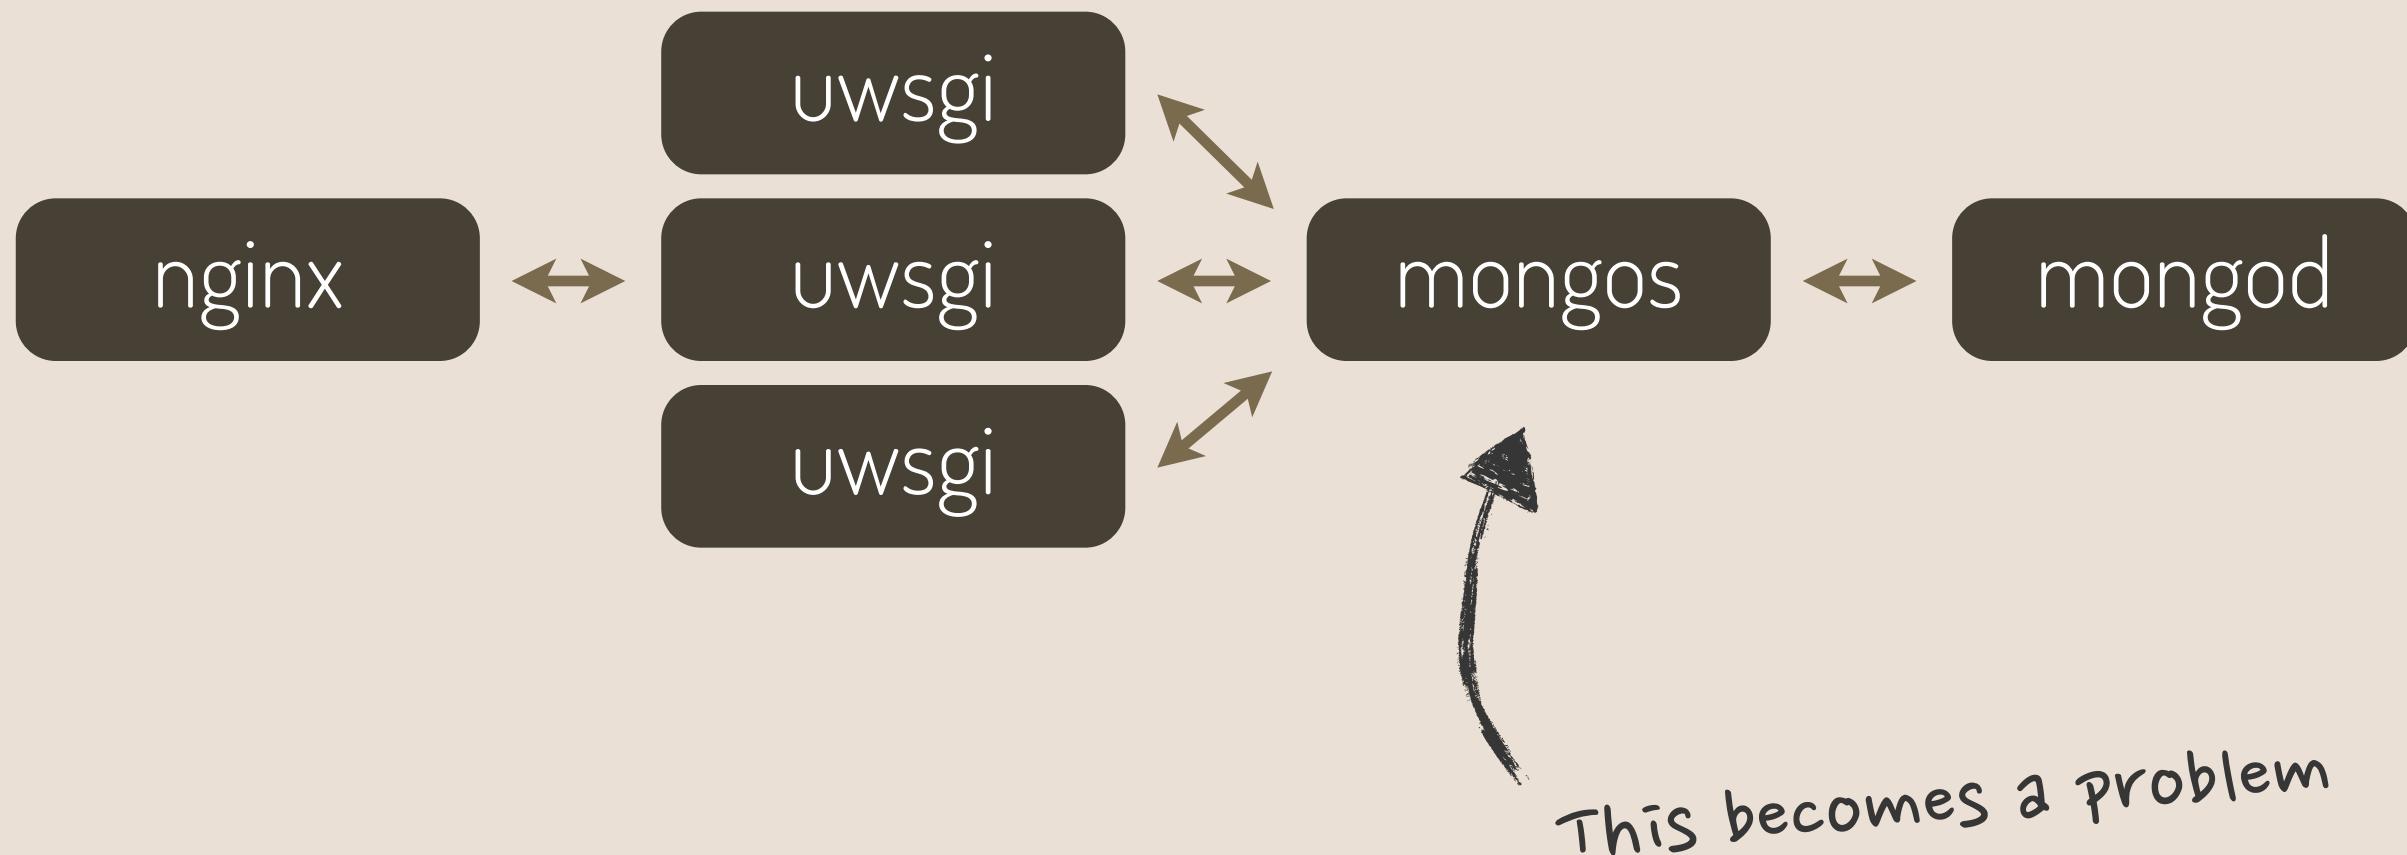
worker setup



worker setup



worker setup



T1 waits for IO

T2 uses CPU

context switch

worker: mongos, give me data

mongos: mongod, give me data

...

mongos: worker, here is your data

context switch

worker: finally! mongos, now give me more data

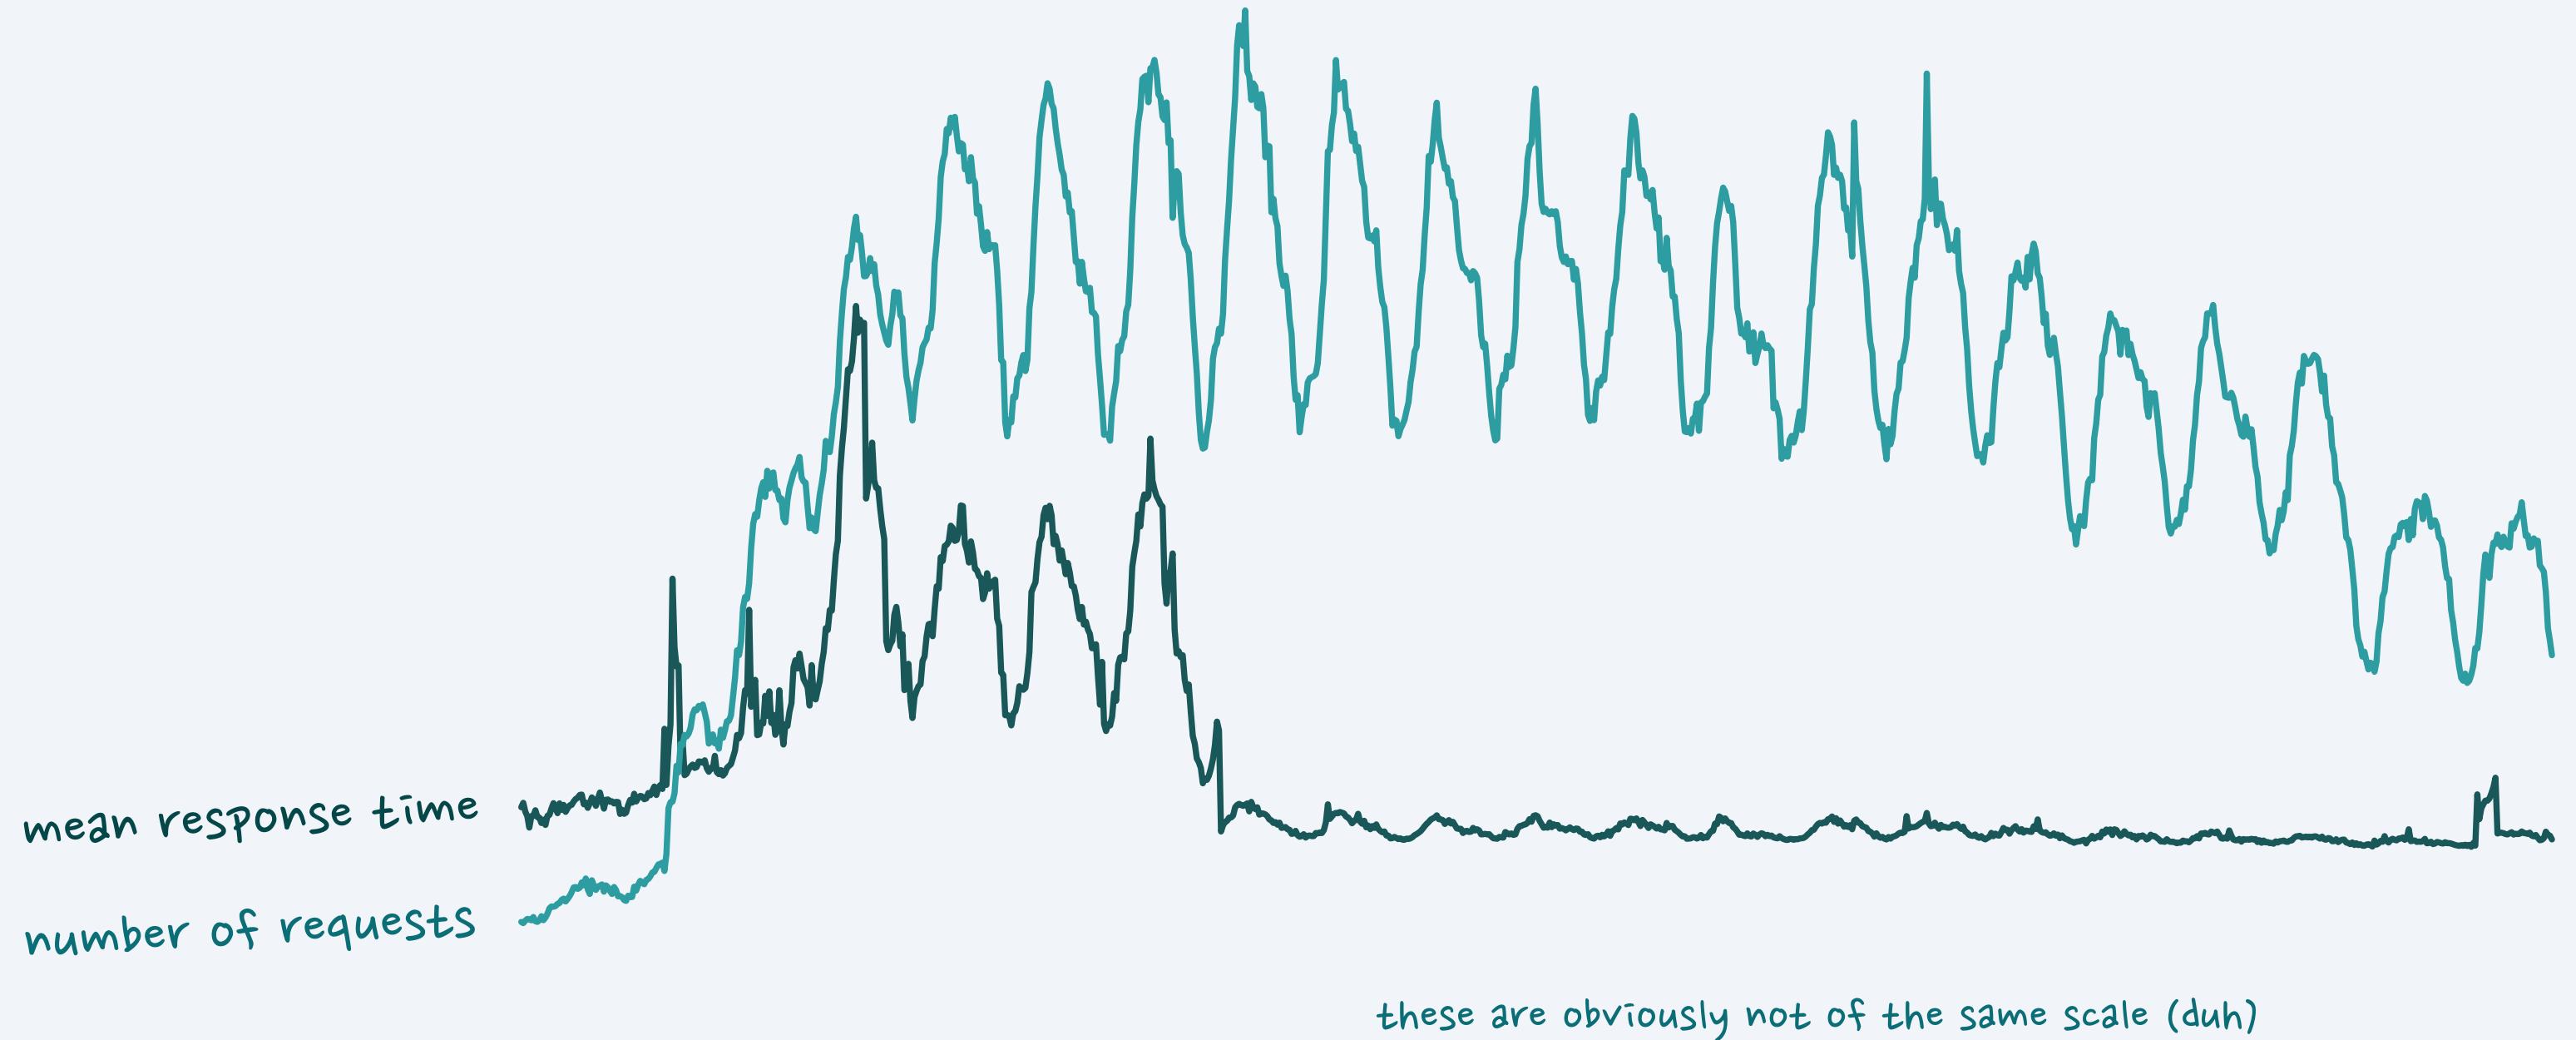
m1.medium: machines with **2 CPUs***
worker and mongos active at the same time
what a novel idea

* cpus[rand() % cpu_choices], might not be an actual CPU



MOAR cPU PLOX

CPU Changes



EBS

it's pretty bad

Breaking your Instance 101

```
$ dd if=/dev/random of=/var/cache/hah bs=4096 count=1024
```

{ MongoDB's Execution Fails }

No transactions

Document-level Operations

No state

transparent reconnects

NO!

Expectation

- mongos fans out and proxies
- if mongos loses connection worker is good
- voluntary primary election is transparent for worker

Actual Result

- mongos fans out well; technically it's a proxy
- if mongos loses connection it terminates both sides
- voluntary primary election kills all connections

MongoDB is Stateful

Tail-able Cursors

`getLastError()`

SIGSEGV

Replica Set Annoyances

1. Add Hidden Secondary
2. Witness it synchronizing
3. Take an existing secondary out
4. Actually unregister the secondary
5. Watch the whole cluster re-elect the same primary
and kill all active connections

Breaking your Cluster 101

- add new primary
- remove old primary
- don't shutdown old primary
- network partitions and one of them overrides the config of the other in the mongoc

{ MongoDB's Design Fails }

Schemaless

Schema vs Schema-less is just a different version of
dynamic typing vs. static typing

Ever since C# and TypeScript:
static typing with an escape hatch to dynamic typing wins

we built an ADT based type system anyways

```
from fireline.schema import types

username = types.String()
profile = types.Dynamic()

x = username.convert('mitsuhiko')
y = profile.convert({'__binary': 'deadbeaf'})
```

GetLastError()

write oddity

why do I need an extra network roundtrip?

write request



mongodb

GetLastError()



mongodb

performance fun

```
import os
from pymongo import Connection

safe = os.environ.get('MONGO_SAFE') == '1'
con = Connection()
db = con['wtfmongo']
coll = db['test']
coll.remove()

for x in xrange(50000):
    coll.insert({'foo': 'bar'}, safe=safe)
```

Disappointing

```
$ MONGO_SAFE=0 time python test.py  
 1.92 real          1.37 user          0.27 sys
```

```
$ MONGO_SAFE=1 time python test.py  
 5.57 real          2.50 user          0.62 sys
```

Disappointing

```
$ MONGO_SAFE=0 time python test.py  
 1.92 real          1.37 user          0.27 sys
```

```
$ MONGO_SAFE=1 time python test.py  
 5.57 real          2.50 user          0.62 sys
```



And that's localhost ...

that would not be a problem if safe mode was fast.
As it stands currently safe mode is slower than Postgres

Lack of Joins

(the shitty map reduce is no replacement)

They will happen

1. Before we had joins, we did not have joins
2. not having joins is not a feature
3. I see people joining in their code by hand. Inefficient

RethinkDB has Distributed Joins :-)

```
r \
  .table('marvel') \
  .inner_join(r.table('dc'),
    Lambda m, dc: m['strength'] < dc['strength']) \
  .run(conn)
```

MongoDB does **not** have Map-Reduce
(that shitty JavaScript map-reduce thing does not count)

Inconsistent Queries

(and a downright dangerous aggregation query system)

Oh got why!?

```
db.bios.find( {  
    "awards": { "$elemMatch": {  
        "award": "Turing Award",  
        "year": { "$gt": 1980 }  
    } }  
} )  
  
db.users.find( { "username": "mitsuhiko" } )
```

Repeat after me: in-band signalling is **wrong!**

Aggregation Framework comes with SQL Injection

```
db.zipcodes.aggregate( {  
    "$group": { "_id": "$state",  
                "total_pop": { "$sum": "$pop" } }  
} , {  
    "$match": { "total_pop": { "$gte": 10 * 1000 * 1000 } }  
} )
```

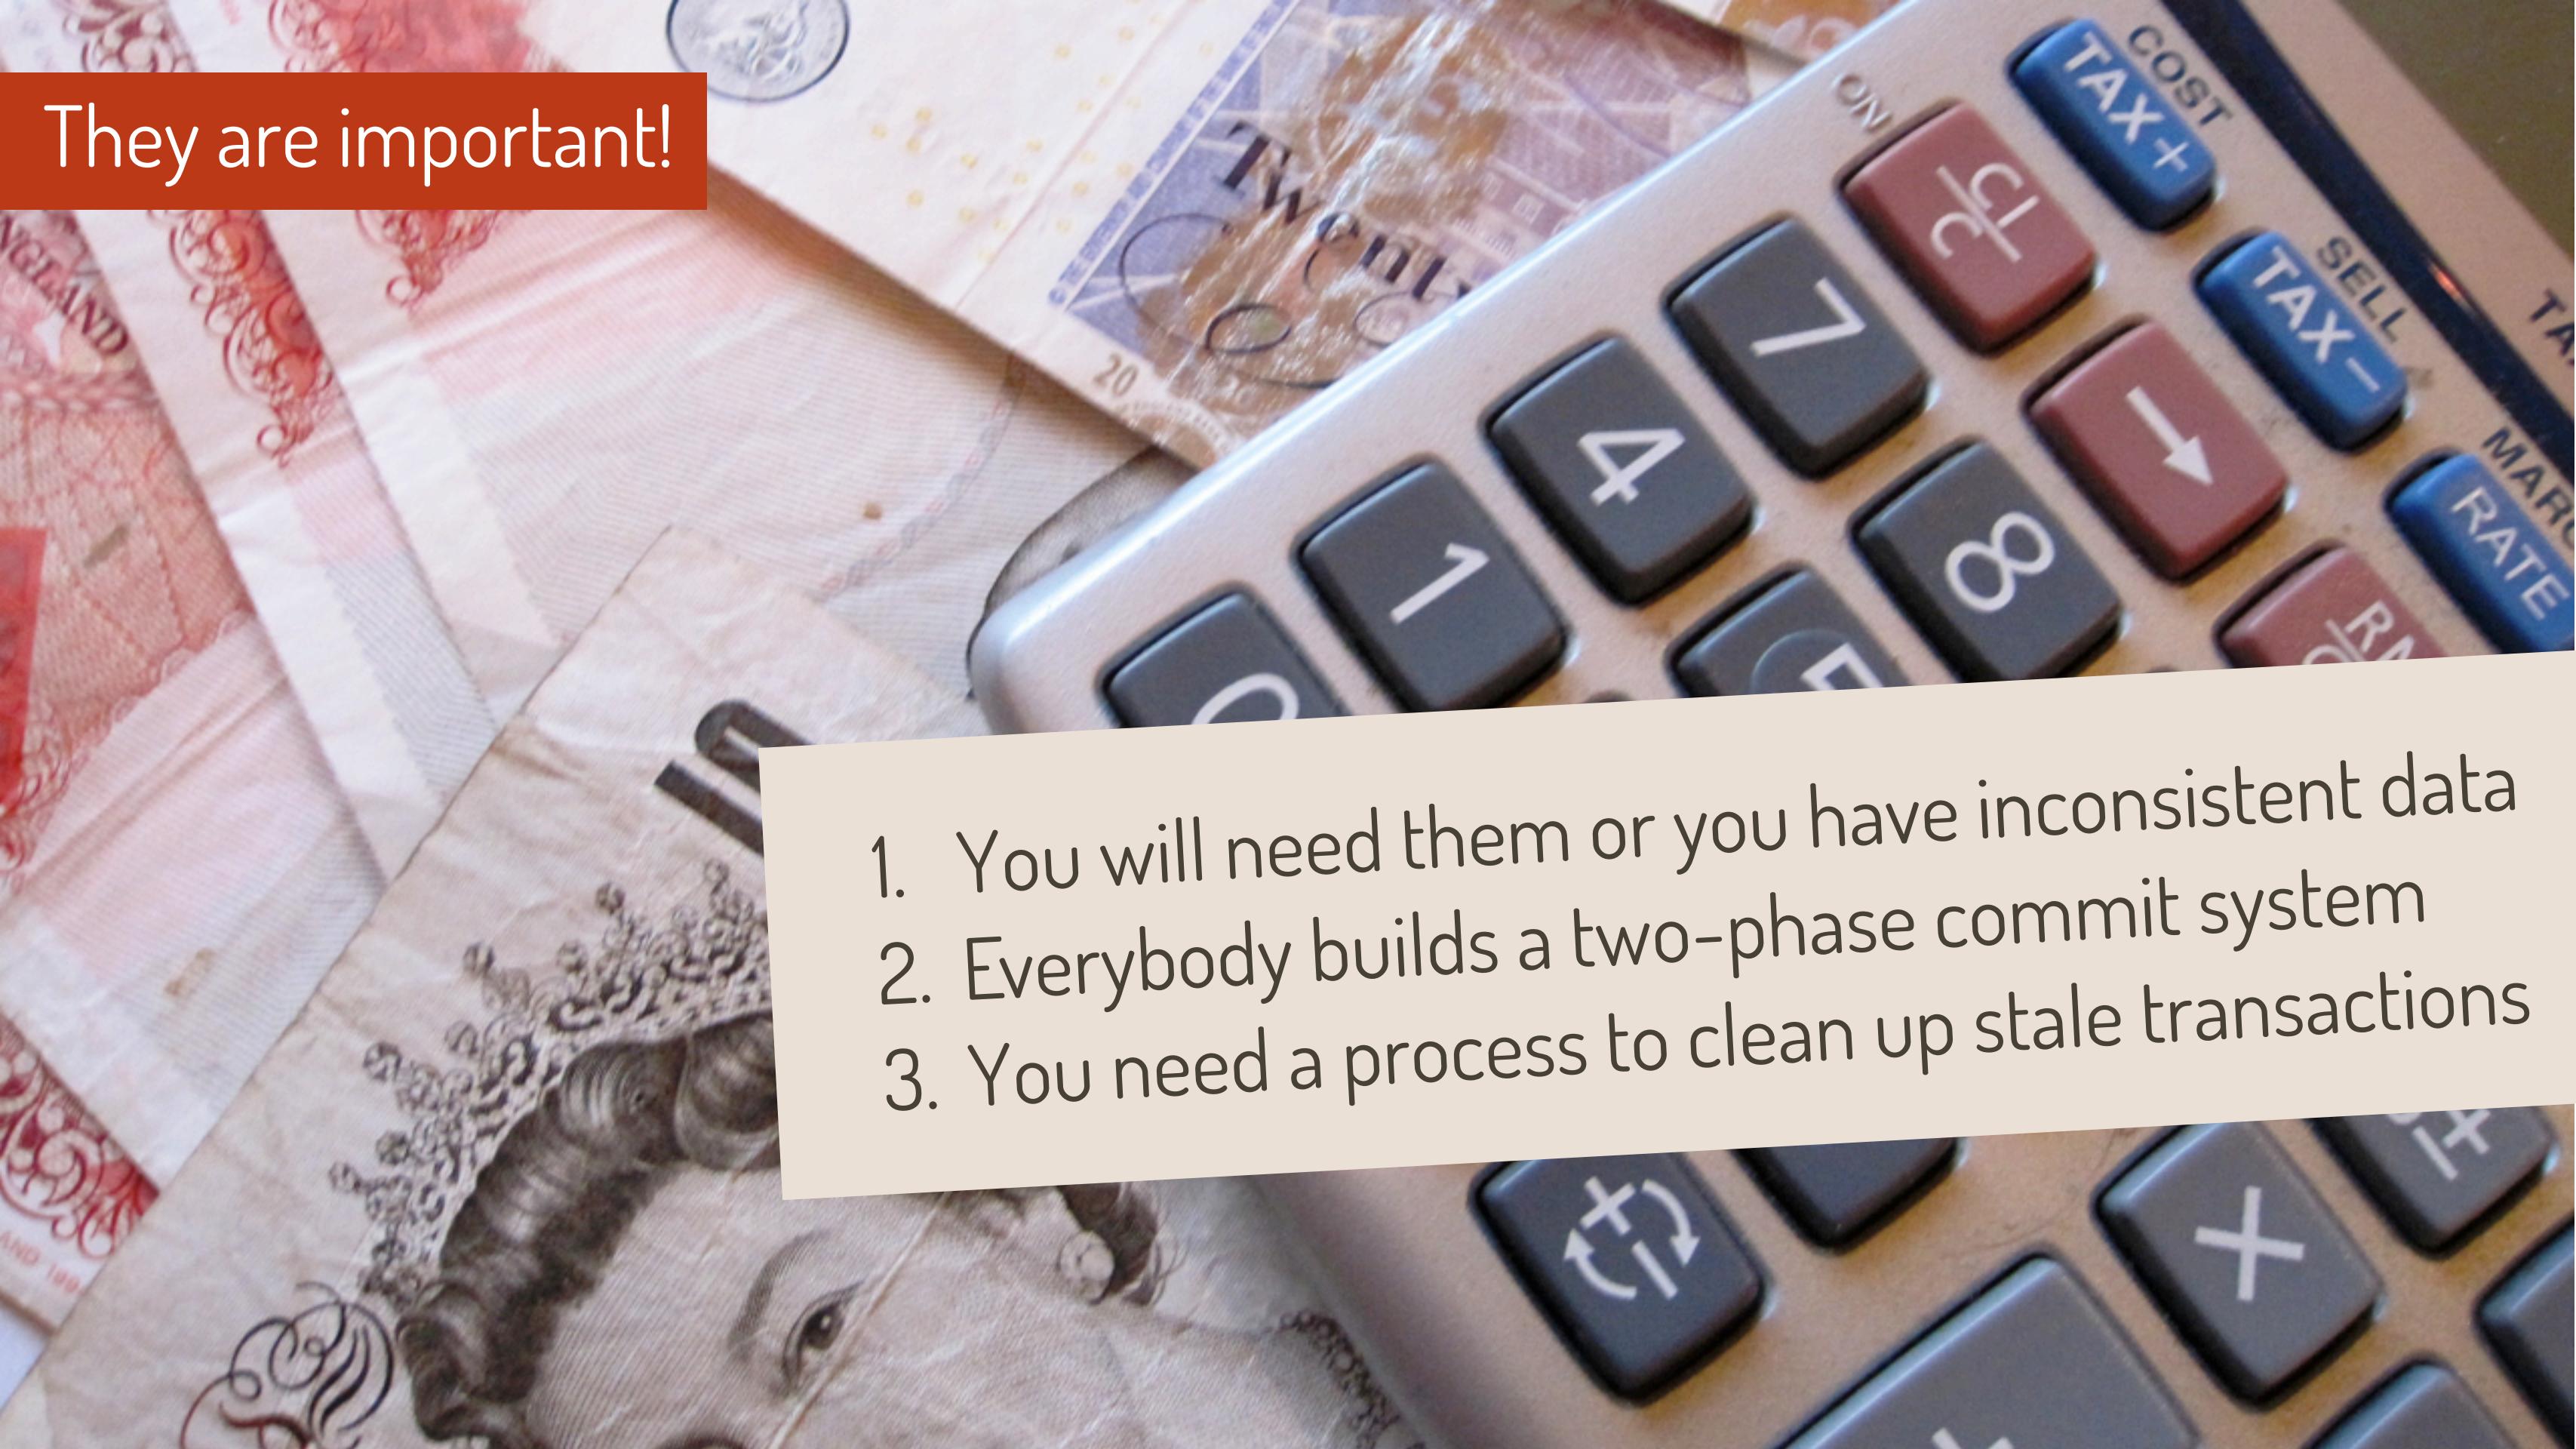
Aggregation Framework comes with SQL Injection

```
db.zipcodes.aggregate( {  
    "$group": { "_id": "$state",  
                "total_pop": { "$sum": "$pop" } }  
} , {  
    "$match": { "total_pop": { "$gte": 10 * 1000 * 1000 } }  
} )
```



spot the injection :-()

No Transactions

A collage of British banknotes (20 and 50 pound) and a calculator. The banknotes are fanned out at the bottom left. A silver electronic calculator is positioned diagonally across the top right, showing various financial functions like 'COST TAX +', 'SELL TAX -', 'MARK RATE', and 'ON'.

They are important!

1. You will need them or you have inconsistent data
2. Everybody builds a two-phase commit system
3. You need a process to clean up stale transactions

Locks Everywhere

A dense pile of numerous padlocks of various colors (blue, red, gold, silver, black) and sizes, all attached to a single metal structure. The padlocks are piled high, creating a textured, metallic surface. Some have small tags or inscriptions on them.

MVCC is good for you

RethinkDB, Postgres and even MySQL support MVCC

Shitty Index Selection

1. MongoDB picks secondary indexes automatically
2. It will also start using sparse indexes
3. It might not give you results back
4. Sometimes forcing ordering makes MongoDB use a compound index

Limited Indexes

1. Given a compound index on [a, b]
2. {a: 1, b: 2} and {\$and: [{a: 1}, {b: 2}]} are equivalent
3. Only the former picks up the compound index
4. Negations never use indexes
5. {\$or: [...]}) is implemented as two parallel queries,
both clauses might need separate indexes.

We have a query optimizer :P

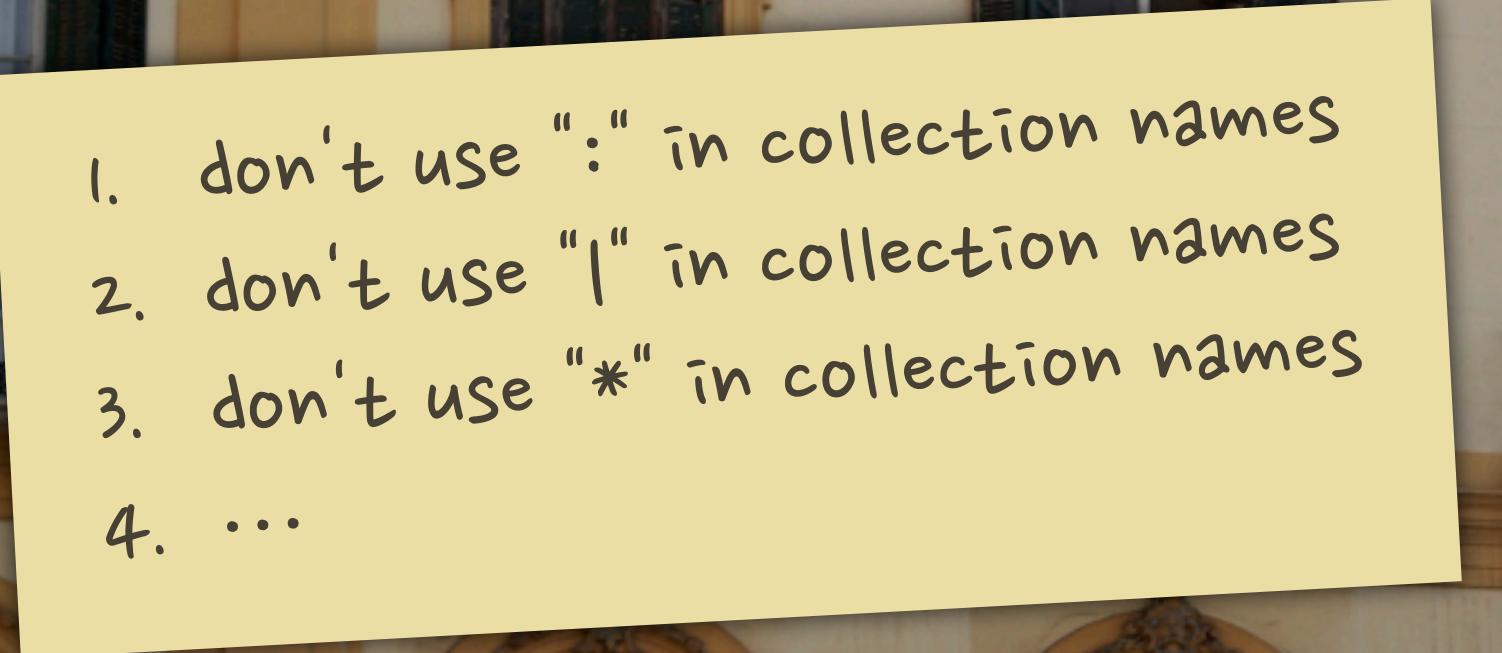
{ Other Things of Note }

Making Mongo not Suck (as much) on OS X

```
$ mongod --noprealloc --smallfiles --nojournal run
```

what are sparse files?

Windows

- 
1. don't use ":" in collection names
 2. don't use "l" in collection names
 3. don't use "*" in collection names
 4. ...



A MongoDB Cluster needs to boot in a certain Order

(Great fun if you have a suspended test infrastructure on Amazon)

MongoDB is a pretty good data dump thing

MongoDB is a pretty good data dump thing
it's not a SQL database

MongoDB is a pretty good data dump thing
it's not a SQL database
but you probably want a SQL database

MongoDB is a pretty good data dump thing
it's not a SQL database
but you probably want a SQL database
at least until RethinkDB is ready

That's it.

Now ask questions.

And add me on twitter: @mitsuhiko

Slides at lucumr.pocoo.org/talks



Legal Shenanigans

Creative Common Sources for Images:

CPU by EssjazNZ: <http://www.flickr.com/photos/essjay/497287571/>

Locks by katiejean97: <http://www.flickr.com/photos/katiejean97/7036715845/>

Money Money Money by Images_of_Money: <http://www.flickr.com/photos/59937401@N07/547416841/>

Through any Window by Josep Ma. Rosell: <http://www.flickr.com/photos/batega/1354354592/in/photostream/>

RAD Soldiers is a Trademark of WarChest Limited. RAD Soldiers Artwork and Logo Copyright © 2013 by WarChest Limited.
All Rights Reserved.